

ASSIGNMENT 1

Assignment Date	25 September 2022
Student Name	Sri Vijaya Harini.R
Student Roll Number	820419205057
Maximum Marks	2 Marks

1. Download the dataset: Dataset

2 . Load the Dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv("gdrive/My Drive/Churn_Modelling.csv")
df.head()
```

OUTPUT:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	4
	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1

```

1      112542.58      0
2      113931.57      1
3      93826.63       0
4      79084.10       0

```

```
from google.colab import drive drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
#dropping row number columns as we already have index column by default
dataset.drop(['RowNumber'], axis=1,inplace=True)
```

3 . Visualizations


```
import matplotlib.pyplot as plt import
seaborn as sns
```

##Univariate Analysis

```
# plt.scatter(churn.index,churn["Age"])
# plt.show()
```


```
# Age Histogram sns.histplot(x='Age',
data=dataset)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f76872b9410>




```
# Geography Histogram sns.histplot(x='Geography',  
data=dataset)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f76864b6390>
```



```
# Geography Histogram sns.histplot(x='Gender',  
data=dataset)
```

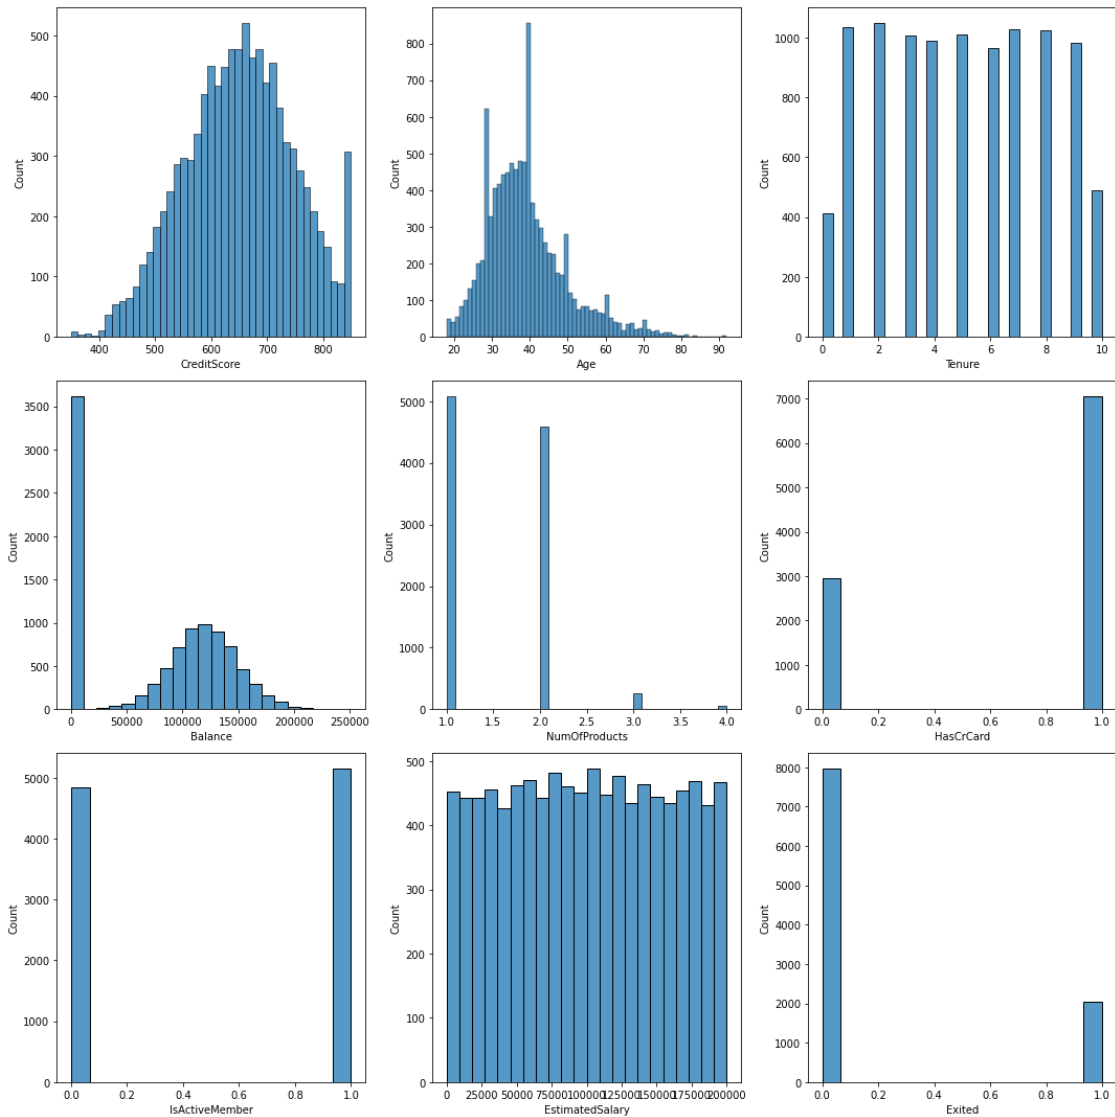
```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7685fdee90>
```



```
cols = 3 rows = 3
num_cols = dataset.select_dtypes(exclude='object').columns #exclude string based columns namely Surname, Geography, Gender print(num_cols)
fig = plt.figure(figsize=(cols*5, rows*5))
for i, col in enumerate(num_cols[1:]): #exclude Customer ID
    ax=fig.add_subplot(rows,cols,i+1)
    sns.histplot(x = dataset[col], ax = ax)


fig.tight_layout() plt.show()

Index(['CustomerId', 'CreditScore', 'Age', 'Tenure', 'Balance',
      'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
      'Exited'], dtype='object')
```




```
# sns.kdeplot(x='Age', data=churn, hue='Exited') sns.histplot(x='Age', data=dataset, kde=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7685ba8290>
```




```
# sns.kdeplot(x='Age', data=churn, hue='IsActiveMember')
sns.histplot(x='CreditScore', data=dataset,
kde=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f768597f2d0>




```
sns.boxplot(x=dataset['Balance'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f7686032110>




```
dataset['Gender'].value_counts().plot.bar()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f7682e1ea50>




```
dataset['Geography'].value_counts().plot.bar()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f7683120d90>



##Bi - Variate Analysis

```
g = sns.pairplot(dataset, diag_kind="kde", markers="+",  
plot_kws=dict(s=50, edgecolor="b", linewidth=1),  
diag_kws=dict(shade=True))
```

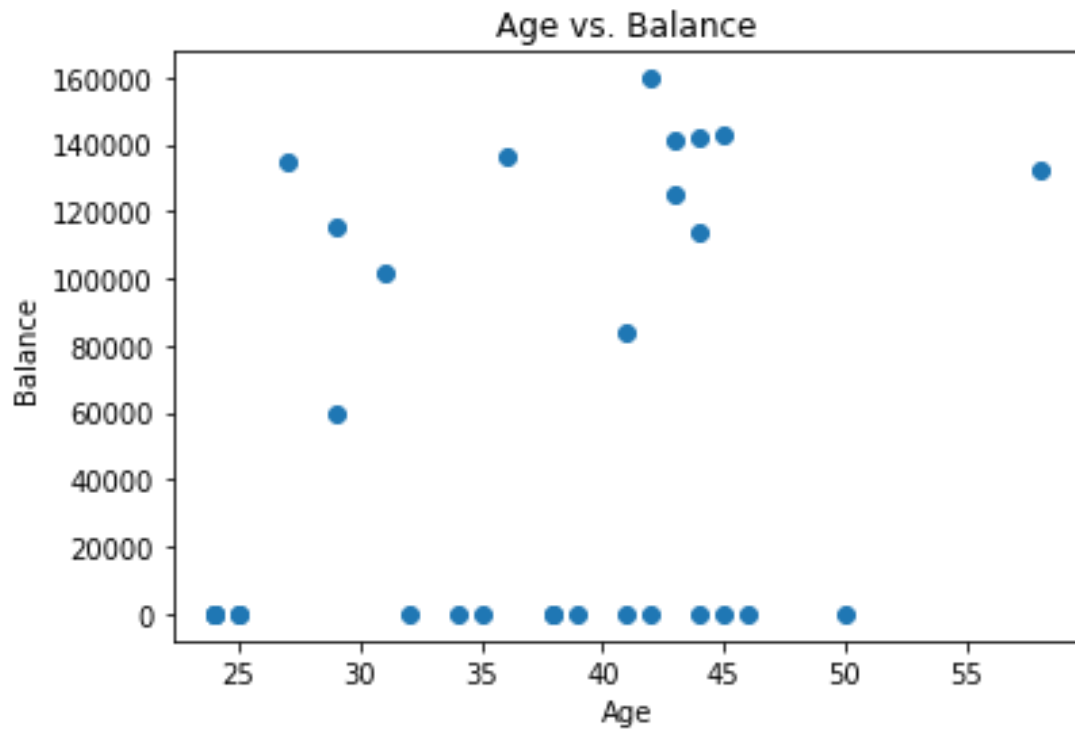



```
import matplotlib.pyplot as plt
```

```
#create scatterplot of hours vs. score
```


```
plt.scatter(dataset.Age[:30], dataset.Balance[:30]) plt.title('Age  
vs. Balance') plt.xlabel('Age') plt.ylabel('Balance')
```

```
Text(0, 0.5, 'Balance')
```



```
dataset.plot.hexbin(x='Age', y='CreditScore', gridsize=10)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7682d84690>
```



```
##Multi-variate Analysis dataset.corr()
```

```

      CustomerId  CreditScore   Age  Tenure  Balance \
CustomerId    1.000000   0.005308  0.009497 -0.014883 -0.012419

```

CreditScore	0.005308	1.000000	-0.003965	0.000842	0.006268
Age	0.009497	-0.003965	1.000000	-0.009997	0.028308
Tenure	-0.014883	0.000842	-0.009997	1.000000	-0.012254
Balance	-0.012419	0.006268	0.028308	-0.012254	1.000000
NumOfProducts	0.016972	0.012238	-0.030680	0.013444	-0.304180
HasCrCard	-0.014025	-0.005458	-0.011721	0.022583	-0.014858
IsActiveMember	0.001665	0.025651	0.085472	-0.028362	-0.010084
EstimatedSalary	0.015271	-0.001384	-0.007201	0.007784	0.012797
Exited	-0.006248	-0.027094	0.285323	-0.014001	0.118533

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	\
CustomerId	0.016972	-0.014025		0.001665	
0.015271					
CreditScore	0.012238	-0.005458		0.025651	-
0.001384					
Age	-0.030680	-0.011721		0.085472	-
0.007201					
Tenure	0.013444	0.022583		-0.028362	
0.007784					
Balance	-0.304180	-0.014858		-0.010084	
0.012797					
NumOfProducts	1.000000	0.003183		0.009612	
0.014204					
HasCrCard	0.003183	1.000000		-0.011866	-
0.009933					
IsActiveMember	0.009612	-0.011866		1.000000	-
0.011421					
EstimatedSalary	0.014204	-0.009933		-0.011421	
1.000000	Exited	-0.047820	-0.007138		-0.156128
0.012097					

	Exited
CustomerId	-0.006248
CreditScore	-0.027094
Age	0.285323
Tenure	-0.014001
Balance	0.118533
NumOfProducts	-0.047820
HasCrCard	-0.007138
IsActiveMember	-0.156128
EstimatedSalary	0.012097
Exited	1.000000

```
sns.set(font_scale=0.50) plt.figure(figsize=(8,4))
sns.heatmap(dataset.corr(),cmap='RdBu_r', annot=True, vmin=-1, vmax=1)
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f7680979950>



#Three variables - Multivariate

```
sns.barplot(x='Age', y='Geography', data=dataset, palette='bright',hue='Gender')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f767ec905d0>



4 . Descriptive statistics import

statistics as st

```
dataset[['Age', 'Balance', 'EstimatedSalary']].mean()
```

```
Age          38.921800
Balance      76485.889288
EstimatedSalary 100090.239881
dtype: float64 dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999 Data
```

```
columns (total 13 columns):
```

```
# Column      Non-Null Count  Dtype
```

```
---  ---
0   CustomerId      10000 non-null  int64
1   Surname         10000 non-null  object
2   CreditScore      10000 non-null  int64
3   Geography       10000 non-null  object
4   Gender          10000 non-null  object
5   Age             10000 non-null  int64
6   Tenure          10000 non-null  int64
7   Balance         10000 non-null  float64
8   NumOfProducts   10000 non-null  int64
9   HasCrCard       10000 non-null  int64
10  IsActiveMember  10000 non-null  int64
11  EstimatedSalary  10000 non-null  float64  12  Exited
      10000 non-null  int64  dtypes: float64(2), int64(8), object(3)
      memory usage: 1015.8+ KB dataset.describe()
```

	CustomerId	CreditScore	Age	Tenure
Balance				
\				
count	1.000000e+04	10000.000000	10000.000000	10000.000000
10000.000000	mean	1.569094e+07	650.528800	38.921800
5.012800	76485.889288	std	7.193619e+04	96.653299
2.892174	62397.405202	min	1.556570e+07	350.000000
0.000000	0.000000	25%	1.562853e+07	584.000000
3.000000	0.000000			32.000000
50%	1.569074e+07	652.000000	37.000000	5.000000
97198.540000	75%	1.575323e+07	718.000000	44.000000
7.000000	127644.240000	max	1.581569e+07	850.000000
92.000000	10.000000	250898.090000		

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
\ count	10000.000000	10000.00000	10000.000000	
10000.000000	mean	1.530200	0.70550	0.515100
100090.239881	std	0.581654	0.45584	0.499797
57510.492818	min	1.000000	0.00000	0.000000
11.580000	25%	1.000000	0.00000	0.000000
51002.110000				
50%	1.000000	1.00000	1.000000	100193.915000
75%	2.000000	1.00000	1.000000	149388.247500
max	4.000000	1.00000	1.000000	199992.480000

```

Exited
count 10000.000000 mean
0.203700 std 0.402769
min 0.000000 25%
0.000000
50% 0.000000 75%
0.000000 max 1.000000
dataset['Age'].median()

37.0

standard_deviation = dataset['CreditScore'].std()
print(standard_deviation) 96.65329873613035
st.mode(dataset['Geography'])

{"type":"string"} st.median(dataset['Age'])

37.0 st.variance(dataset['CreditScore'])

9341.860156575658

```

5 . Handle Missing Values `dataset.isnull().sum()`

#no missing values

CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0

```
NumOfProducts    0
HasCrCard         0
IsActiveMember    0
EstimatedSalary   0
Exited            0 dtype: int64
```


6 . Find and replace outliers

Visualize Outliers `sns.boxplot(dataset['CreditScore'],data=dataset)`

`/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:`


FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation. FutureWarning

`<matplotlib.axes._subplots.AxesSubplot at 0x7f767ebd9d90>`



`dataset['Balance'].hist()`

`<matplotlib.axes._subplots.AxesSubplot at 0x7f767ebbefd0>`



```
for col in num_cols[1:]:
    print('skewness value of ',col,dataset[col].skew())
```

#Skewness should be in the range of -1 to 1, any columns with skewness outside of that range would have outliers

```
skewness value of  CreditScore -0.07160660820092675
skewness value of  Age 1.0113202630234552 skewness
value of  Tenure 0.01099145797717904 skewness value of
Balance -0.14110871094154384 skewness value of
NumOfProducts 0.7455678882823168 skewness value of
HasCrCard -0.9018115952400578 skewness value of
IsActiveMember -0.06043662833499078 skewness value of
EstimatedSalary 0.0020853576615585162 skewness value of
Exited 1.4716106649378211
```

```
Q1=dataset['Age'].quantile(0.25)
Q3=dataset['Age'].quantile(0.75)
IQR=Q3-Q1
```

```
IQR
```

```
12.0
```

Removing Outliers

#Values above than the upper bound and below than the lower bound are considered outliers

```
upper = dataset['Age'] >= (Q3+1.5*IQR)
```



```
# print("Upper bound:",upper) print(np.where(upper))
```

```
lower = dataset['Age'] <= (Q1-1.5*IQR)
# print("Lower bound:", lower)
print(np.where(lower))
```

```
(array([ 58,  85, 104, 158, 181, 230, 234, 243, 252, 276, 310,
        364, 371, 385, 387, 399, 538, 559, 567, 602, 612, 617,
        658, 678, 696, 736, 766, 769, 807, 811, 823, 859, 884,
        888, 948, 952, 957, 963, 969, 997, 1009, 1039, 1040, 1055,
        1114, 1205, 1234, 1235, 1246, 1252, 1278, 1285, 1328, 1342, 1387,
        1407, 1410, 1433, 1439, 1457, 1519, 1543, 1607, 1614, 1642, 1790,
        1810, 1866, 1901, 1904, 1907, 1933, 1981, 1996, 2002, 2012, 2039,
        2053, 2078, 2094, 2108, 2154, 2159, 2164, 2244, 2274, 2433, 2458,
        2459, 2519, 2553, 2599, 2615, 2659, 2670, 2713, 2717, 2760, 2772,
        2778, 2791, 2855, 2877, 2901, 2908, 2925, 2926, 3008, 3033, 3054,
        3110, 3142, 3166, 3192, 3203, 3229, 3305, 3308, 3311, 3314, 3317,
        3346, 3366, 3368, 3378, 3382, 3384, 3387, 3396, 3403, 3434, 3462,
        3497, 3499, 3527, 3531, 3541, 3559, 3573, 3575, 3593, 3602, 3641,
        3646, 3647, 3651, 3690, 3691, 3702, 3719, 3728, 3733, 3761, 3813,
        3826, 3880, 3881, 3888, 3909, 3910, 3927, 3940, 3980, 3994, 4010,
        4025, 4048, 4051, 4095, 4142, 4147, 4157, 4162, 4170, 4241, 4244,
        4256, 4273, 4280, 4297, 4313, 4318, 4335, 4360, 4366, 4378, 4387,
        4396, 4435, 4438, 4463, 4490, 4501, 4506, 4559, 4563, 4590, 4595,
        4644, 4698, 4747, 4751, 4801, 4815, 4832, 4849, 4931, 4947, 4966,
        4992, 5000, 5020, 5038, 5068, 5132, 5136, 5148, 5159, 5197, 5223,
        5225, 5235, 5255, 5299, 5313, 5368, 5377, 5405, 5457, 5490, 5508,
        5514, 5576, 5577, 5581, 5655, 5660, 5664, 5671, 5698, 5777, 5783,
        5817, 5825, 5840, 5867, 5907, 5957, 5996, 6046, 6116, 6152, 6166,
        6167, 6173, 6212, 6230, 6278, 6289, 6315, 6357, 6366, 6373, 6375,
        6410, 6443, 6515, 6530, 6532, 6581, 6612, 6626, 6706, 6709, 6715,
        6721, 6759, 6763, 6812, 6899, 6970, 6997, 7008, 7057, 7058, 7063,
        7071, 7078, 7094, 7138, 7139, 7142, 7156, 7194, 7202, 7238, 7243,
        7272, 7302, 7362, 7375, 7392, 7499, 7514, 7523, 7526, 7548, 7552,
        7624, 7629, 7692, 7694, 7709, 7715, 7719, 7720, 7727, 7773, 7776,
        7784, 7788, 7802, 7813, 7851, 7894, 7898, 7933, 7956, 7995, 8019,
        8037, 8094, 8098, 8156, 8193, 8207, 8217, 8304, 8321, 8385, 8394,
        8444, 8458, 8467, 8469, 8478, 8488, 8562, 8568, 8577, 8602, 8674,
        8686, 8689, 8711, 8759, 8761, 8768, 8787, 8793, 8822, 8865, 8900,
        8917, 8930, 9018, 9062, 9080, 9112, 9116, 9162, 9223, 9279, 9292,
        9309, 9318, 9324, 9332, 9333, 9351, 9380, 9402, 9425, 9428, 9438,
        9472, 9490, 9506, 9555, 9557, 9582, 9587, 9589, 9593, 9646, 9671,
        9673, 9681, 9686, 9688, 9718, 9733, 9734, 9736, 9747, 9753, 9765,
        9832, 9879, 9894, 9936]),) (array([],
dtype=int64),)
```

```
#Removing outliers based off Age column
```

```
# IQR
```

```
Q1 = np.percentile(dataset['Age'], 25,  
interpolation = 'midpoint')
```

```
Q3 = np.percentile(dataset['Age'], 75,  
interpolation = 'midpoint')
```

```
IQR = Q3 - Q1
```

```
print("Old Shape: ", dataset.shape)
```

```
# Upper bound upper = np.where(dataset['Age'] >=  
(Q3+1.5*IQR))
```

```
# Lower bound lower = np.where(dataset['Age'] <=  
(Q1-1.5*IQR))
```

```
''' Removing the Outliers ''' dataset.drop(upper[0],  
inplace = True) dataset.drop(lower[0], inplace =  
True)
```

```
print("New Shape: ", dataset.shape)
```

```
Old Shape: (10000, 13) New
```

```
Shape: (9589, 13) dataset
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	\
0	15634602	Hargrave	619	France	Female	42	2	
1	15647311	Hill	608	Spain	Female	41	1	
2	15619304	Onio	502	France	Female	42	8	
3	15701354	Boni	699	France	Female	39	1	
4	15737888	Mitchell	850	Spain	Female	43	2	
...	
9995	15606229	Obijiaku	771	France	Male	39	5	
9996	15569892	Johnstone	516	France	Male	35	10	
9997	15584532	Liu	709	France	Female	36	7	
9998	15682355	Sabbatini	772	Germany	Male	42	3	
9999	15628319	Walker	792	France	Female	28		
4								

	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	\
0	0.00	1	1	1	101348.88	
1	83807.86		1	0	1	
	112542.58					
2	159660.80		3	1	0	
	113931.57					
3	0.00	2	0	0	93826.63	

4	125510.82		1	1		1	
	79084.10	
				
9995	0.00		2	1		0	96270.64
9996	57369.61		1		1		
	101699.77					1	
9997	0.00		1		0		42085.58
9998	75075.31		2		1		
	92888.52					0	
9999	130142.79		1		1		
	38190.78		Exited			0	
0	1						
1	0						
2	1						
3	0						
4	0				
9995	0						
9996	0						
9997	1						
9998	1						
9999	0						

[9589 rows x 13 columns]

```
for col in num_cols[1:]:
    print('skewness value of ',col,dataset[col].skew())
```

*# Now we have reduced the Age column's skewness values within -1 to 1 range #
We left the Exited column's skewness value as it is the dependent variable*

```
skewness value of  CreditScore -0.07274225895185718
skewness value of  Age 0.44721544739487257 skewness value
of  Tenure 0.008085830714996462 skewness value of
Balance -0.1409005824644143 skewness value of
NumOfProducts 0.7470530176747141 skewness value of
HasCrCard -0.9034483996482451 skewness value of
IsActiveMember -0.008552881368996219 skewness value of
EstimatedSalary -0.0025661797132480266 skewness value of
Exited 1.4798502461410206
```

7 . Check for Categorical columns and perform encoding

```
##Label encoding and One Hot encoding dataset.reset_index(inplace=True)
```

```
from sklearn.preprocessing import LabelEncoder from
sklearn.preprocessing import OneHotEncoder from
sklearn.compose import ColumnTransformer
```

```
categorical_feature_mask = dataset.dtypes==object
categorical_cols = dataset.columns[categorical_feature_mask].tolist()
```

```
categorical_cols=categorical_cols[1:]
categorical_cols
```

```
['Geography', 'Gender']
```

```
le = LabelEncoder()
dataset[categorical_cols] =
```

```
dataset[categorical_cols].apply(lambda col:
```

```
le.fit_transform(col))
dataset[categorical_cols].head(10)
```

	Geography	Gender
0	0	0
1	2	0
2	0	0
3	0	0
4	2	0
5	2	1
6	0	1
7	1	0
8	0	1

```
categorical_feature_mask
```

```
index      False CustomerId
```

```
False
```

```
Surname      True
```

```
CreditScore  False
```

```
Geography    True
```

```
Gender       True
```

```
Age          False
```

```
Tenure       False
```

```
Balance      False
```

```
NumOfProducts False
```

```
HasCrCard    False
```

```
IsActiveMember False
```

```
EstimatedSalary False
```

```
Exited      False dtype:
```

```
bool
```

```
enc=OneHotEncoder()
```

```
enc_data=pd.DataFrame(enc.fit_transform(dataset[['Geography', 'Gender']]).toarray())
enc_data
```

	0	1	2	3	4
0	1.0	0.0	0.0	1.0	0.0
1	0.0	0.0	1.0	1.0	0.0
2	1.0	0.0	0.0	1.0	0.0
3	1.0	0.0	0.0	1.0	0.0
4	0.0	0.0	1.0	1.0	0.0

```

9584  1.0  0.0  0.0  0.0  1.0
9585  1.0  0.0  0.0  0.0  1.0
9586  1.0  0.0  0.0  1.0  0.0
9587  0.0  1.0  0.0  0.0  1.0
9588  1.0  0.0  0.0  1.0  0.0

```

[9589 rows x 5 columns]

#First three columns of enc_data is for Geography and the next two columns is for Gender, we can replace the already existing categorical columns with these encoded values

#Dropping already existing Geography and Gender columns

```

dataset.drop(['Geography'], axis=1,inplace=True)
dataset.drop(['Gender'], axis=1,inplace=True)

```

```

dataset.insert(2, "Geography_France", enc_data.iloc[:,0], True)
dataset.insert(3, "Geography_Germany", enc_data.iloc[:,1], True)
dataset.insert(4, "Geography_Spain", enc_data.iloc[:,2], True)
dataset.insert(5, "Gender_Female", enc_data.iloc[:,3], True)
dataset.insert(6, "Gender_Male", enc_data.iloc[:,4], True) dataset

```

	index	CustomerId	Geography_France	Geography_Germany	Geography_Spain
\					
0	0	15634602	1.0	0.0	0.0
1	1	15647311	0.0	0.0	1.0
2	2	15619304	1.0	0.0	0.0
3	3	15701354	1.0	0.0	0.0
4	4	15737888	0.0	0.0	1.0

	...				
9584	9995	15606229	1.0	0.0	0.0
9585	9996	15569892	1.0	0.0	0.0
9586	9997	15584532	1.0	0.0	0.0
9587	9998	15682355	0.0	1.0	0.0
	9588	9999	15628319	1.0	0.0
	0.0				

	Gender_Female	Gender_Male	Surname	CreditScore	Age	Tenure	\
0	1.0	0.0	Hargrave	619	42	2	
1	1.0	0.0	Hill	608	41	1	
2	1.0	0.0	Onio	502	42	8	
3	1.0	0.0	Boni	699	39	1	
4	1.0	0.0	Mitchell	850	43	2	
...	
9584	0.0	1.0	Obijiaku	771	39	5	
9585	0.0	1.0	Johnstone	516	35	10	
9586	1.0	0.0	Liu	709	36	7	

9587	0.0	1.0	Sabbatini	772	42	3	9588
	1.0	0.0	Walker	792	28	4	

	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	\
0	0.00	1	1	1	101348.88	
1	83807.86		1	0	1	
	112542.58					
2	159660.80		3	1	0	
	113931.57					
3	0.00	2	0	0	93826.63	
4	125510.82		1	1	1	
	79084.10	
				
9584	0.00	2	1	0	96270.64	
9585	57369.61		1	1	1	
	101699.77					
9586	0.00	1	0	1	42085.58	
9587	75075.31		2	1	0	
	92888.52	9588	130142.79	1	1	
	0	38190.78				

	Exited
0	1
1	0
2	1
3	0
4	0
	...
9584	0
9585	0
9586	1
9587	1
9588	0

[9589 rows x 17 columns]

We drop some irrelevant columns that does not contribute to prediction

```
dataset.drop(columns="CustomerId",axis=1,inplace=True)
dataset.drop(columns="Surname",axis=1,inplace=True)
dataset.drop(columns="index",axis=1,inplace=True) dataset
```

	Geography_France	Geography_Germany	Geography_Spain	Gender_Female	\
0	1.0	0.0	0.0	1.0	
1	0.0	0.0	1.0	1.0	
2	1.0	0.0	0.0	1.0	
3	1.0	0.0	0.0	1.0	
4	0.0	0.0	1.0	1.0	

	...				

9584	1.0	0.0	0.0	0.0
9585	1.0	0.0	0.0	0.0
9586	1.0	0.0	0.0	1.0
9587	0.0	1.0	0.0	0.0
	9588	1.0	0.0	0.0
	1.0			

	Gender_Male	CreditScore	Age	Tenure	Balance	NumOfProducts	\
0	0.0	619	42	2	0.00	1	
1	0.0	608	41	1	83807.86	1	
2	0.0	502	42	8	159660.80	3	
3	0.0	699	39	1	0.00	2	
4	0.0	850	43	2	125510.82	1	...
	
9584	1.0	771	39	5	0.00	2	
9585	1.0	516	35	10	57369.61	1	
9586	0.0	709	36	7	0.00	1	
9587	1.0	772	42	3	75075.31	2	9588
	0.0	792	28	4	130142.79	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0

9584	1	0	96270.64	0
9585	1	1	101699.77	0
9586	0	1	42085.58	1
9587	1	0	92888.52	1
9588	1	0	38190.78	0

[9589 rows x 14 columns]

8 . Split the data into dependent and independent variables

X= dataset.iloc[:, :-1].values *#Indepedent variables*
y= dataset.iloc[:, -1].values *#Dependent variables*

X

```
array([[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
        1.0000000e+00, 1.0134888e+05],
       [0.0000000e+00, 0.0000000e+00, 1.0000000e+00, ..., 0.0000000e+00,
        1.0000000e+00, 1.1254258e+05],
       [1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
```

```

0.0000000e+00, 1.1393157e+05],
...,
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
1.0000000e+00, 4.2085580e+04],
[0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
0.0000000e+00, 9.2888520e+04],
[1.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
0.0000000e+00, 3.8190780e+04]]) y

array([1, 0, 1, ..., 1, 1, 0])

```

9 . Scale the independent variable

```

from sklearn.preprocessing import StandardScaler
scale= StandardScaler() X
= scale.fit_transform(X)
X

```

```

array([[ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
0.99573337, 0.01997639],
[-1.0028197 , -0.57955796, 1.74527693, ..., -1.54891873,
0.99573337, 0.21465635],
[ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
-1.00428491, 0.23881355],
...,
[ 0.99718823, -0.57955796, -0.57297497, ..., -1.54891873,
0.99573337, -1.01072631],
[-1.0028197 , 1.72545295, -0.57297497, ..., 0.64561166,
-1.00428491, -0.12716553],
[ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
-1.00428491, -1.07846436]])

```

10 . Split the data into training and testing

```

from sklearn.model_selection import train_test_split

```

```

# We use train_test_split function to split the data such that 25% is used
for testing while the remaining 75% is used for training
X_train, X_test, y_train, y_test = train_test_split(X,y ,
random_state=104,test_size=0.25, shuffle=True)

```

```

X_train

```

```

array([[ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
0.99573337, -1.74019169],
[-1.0028197 , -0.57955796, 1.74527693, ..., -1.54891873,
-1.00428491, -1.39787901],
[-1.0028197 , 1.72545295, -0.57297497, ..., -1.54891873,

```



```

    0.99573337, -1.48817335],
    ...,
    [ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
      -1.00428491, 0.71481237],
    [ 0.99718823, -0.57955796, -0.57297497, ..., -1.54891873,
      -1.00428491, 0.60834563],
    [-1.0028197 , 1.72545295, -0.57297497, ..., 0.64561166,
      0.99573337, 0.0525285 ]]) X_test

array([[ -1.0028197 , -0.57955796, 1.74527693, ..., -1.54891873,
        -1.00428491, -0.90389608],
       [ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
         0.99573337, -0.54087223],
       [-1.0028197 , -0.57955796, 1.74527693, ..., 0.64561166,
         0.99573337, -1.02004733],
       ...,
       [ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
         0.99573337, -0.23978536],
       [ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
         0.99573337, -0.17457887],
       [ 0.99718823, -0.57955796, -0.57297497, ..., 0.64561166,
        -1.00428491, -0.0121091 ]]) y_train

array([0, 0, 0, ..., 0, 0, 0]) y_test

array([0, 1, 0, ..., 0, 0, 1])

```